

Problem 1: Linear Regression from Scratch (30 points)

```
In [1]: # import the necessary packages
import numpy as np
from matplotlib import pyplot as plt
np.random.seed(100)
```

Let's generate some data points first, by the equation $y = x - 3$.


```
In [2]: x = np.random.randint(100, size=100)/30 - 2
X = x.reshape(-1, 1)

y = x + -3 + 0.3*np.random.randn(100)
```

Let's then visualize the data points we just created.

```
In [3]: plt.scatter(X, y)
```

```
Out[3]: <matplotlib.collections.PathCollection at 0x7fcd52239be0>
```



1.1 Gradient of vanilla linear regression model (5 points)

In the lecture, we learn that the cost function of a linear regression model can be expressed as **Equation 1**:

$$J(\theta) = \frac{1}{2m} \sum_i^m \left(h_{\theta} \left(x^{(i)} \right) - y^{(i)} \right)^2$$

The gradient of it can be written as **Equation 2**:

$$\frac{\partial J(\theta)}{\partial \theta} = \frac{1}{m} \sum_i^m \left(x^{(i)} \right) \left(h_{\theta} \left(x^{(i)} \right) - y^{(i)} \right)$$

1.2 Gradient of vanilla regularized regression model (5 points)

After adding the L2 regularization term, the linear regression model can be expressed as **Equation 3**:

$$J(\theta) = \frac{1}{2m} \sum_i^m \left(h_{\theta} \left(x^{(i)} \right) - y^{(i)} \right)^2 + \frac{\lambda}{2m} \sum_j^n (\theta_j)^2$$

The gradient of it can be written as **Equation 4**:

$$\frac{\partial J(\theta)}{\partial \theta} = \frac{1}{m} \sum_i^m \left(x^{(i)} \right) \left(h_{\theta} \left(x^{(i)} \right) - y^{(i)} \right) + \frac{\lambda}{m} \sum_j^n (\theta_j)$$

1.3 Implement the cost function of a regularized regression model (5 points)

Please implement the cost function of a regularized regression model according to the above equations.

1.4 Implement the gradient of the cost function of a regularized regression model (5 points)

Please implement the gradient of the cost function of a regularized regression model according to the above equations.

```
In [4]: def regularized_linear_regression(X, y, alpha=0.01, lambda_value=1, epochs=30):
        """
        :param x: feature matrix
        :param y: target vector
        :param alpha: learning rate (default:0.01)
        :param lambda_value: lambda (default:1)
        :param epochs: maximum number of iterations of the
                        linear regression algorithm for a single run (default=30)
        :return: weights, list of the cost function changing overtime
        """

        m = np.shape(X)[0] # total number of samples
        n = np.shape(X)[1] # total number of features

        X = np.concatenate((np.ones((m, 1)), X), axis=1)
        W = np.random.randn(n + 1, )

        # stores the updates on the cost function (loss function)
        cost_history_list = []

        # iterate until the maximum number of epochs
        for current_iteration in np.arange(epochs): # begin the process

            # compute the dot product between our feature 'X' and weight 'W'
            y_estimated = X.dot(W)

            # calculate the difference between the actual and predicted value
            error = y_estimated - y
```

```
#####
##### Begin of Question 1.3 #####
#####
```

```

##### Please write down your code here:#####

# calculate the cost (MSE) (Equation 1)
cost_without_regularization = (1/(2*m))*(error.dot(error))

##### Please write down your code here:#####

# regularization term
reg_term = (lambda_value/(2*m))*(W.dot(W))

# calculate the cost (MSE) + regularization term (Equation 3)
cost_with_regularization = cost_without_regularization + reg_term

#####
##### End of Question 1.3 #####
#####

#####
##### Begin of Question 1.4 #####
#####

##### Please write down your code here:#####

# calculate the gradient of the cost function with regularization term (Equation
gradient = (1/m)*(X.T.dot(error)) + (lambda_value/m)*(W)

# Now we have to update our weights
W = W - alpha * gradient

#####
##### End of Question 1.4 #####
#####

# keep track the cost as it changes in each iteration
cost_history_list.append(cost_with_regularization)

# Let's print out the cost
print(f"Cost with regularization: {cost_with_regularization}")
print(f"Mean square error: {cost_without_regularization}")

return W, cost_history_list

```

Run the following code to train your model.

Hint: If you have the correct code written above, the cost should be 0.5181222986588751 when $\lambda = 10$.

```

In [5]: lambda_list = [0, 10, 100, 1000, 10000]
for lambda_ in lambda_list:
    # calls regression function with different values of lambda
    weight, _ = regularized_linear_regression(X, y, alpha=0.01,
                                              lambda_value=lambda_, epochs=1000)


    fitted_line = np.dot(X, weight[1]) + weight[0]
    plt.scatter(X, y, label='data points')
    plt.plot(X, fitted_line, color='r', label='Fitted line')
    plt.xlabel("X")
    plt.ylabel("Y")

```

```
plt.title(f"Regression (lambda : {lambda_})")
plt.legend()
plt.show()
```


Cost with regularization: 0.05165888565058274

Mean square error: 0.05165888565058274




Cost with regularization: 0.5181225049184746

Mean square error: 0.08982014821513126




Cost with regularization: 2.793172488740026

Mean square error: 1.2785107029715974




Cost with regularization: 5.591464362606628

Mean square error: 4.946888025066496



Cost with regularization: 6.242695626933972
Mean square error: 6.161442583355812



1.5 Analyze your results (10 points)

According to the above figures, what's the best choice of λ ?

Why the regressed line turns to be flat as we increase λ ?

The best choice for λ is 0 because it has the lowest cost with regularization and mean square error.

The regressed line turns to be flat as λ increases because as λ gets larger this makes the coefficients less important and pushes them towards zero.

Problem 2: Getting familiar with PyTorch (30 points)

```
In [6]: import mltools as ml
import torch
```

```
In [7]: data = np.genfromtxt("data/curve80.txt")
X = data[:,0]
X = np.atleast_2d(X).T # code expects shape (M,N) so make sure it's 2-dimensional
Y = data[:,1] # doesn't matter for Y
Xtr,Xte,Ytr,Yte = ml.splitData(X,Y,0.75) # split data set 75/25
```

```
degree = 5
XtrP = ml.transforms.fpoly(Xtr, degree=degree, bias=False)
XtrP, params = ml.transforms.rescale(XtrP)
```

```
In [8]: XtrP_tensor = torch.from_numpy(XtrP)
Ytr_tensor = torch.from_numpy(np.reshape(Ytr, (Ytr.shape[0],1)))

XtrP_tensor = XtrP_tensor.float()
Ytr_tensor = Ytr_tensor.float()
```


```
In [9]: linear_regressor = torch.nn.Linear(in_features=XtrP_tensor.shape[1], out_features=Ytr_te
```

```
In [10]: criterion = torch.nn.MSELoss()
optimizer = torch.optim.SGD(linear_regressor.parameters(), lr=0.1)
epochs = 100000
```

```
In [11]: loss_record = []
for _ in range(epochs):
    optimizer.zero_grad() # set gradient to zero
    pred_y = linear_regressor(XtrP_tensor)
    loss = criterion(pred_y, Ytr_tensor) # calculate loss function
    loss.backward() # backpropagate gradient
    loss_record.append(loss.item())
    optimizer.step() # update the parameters in the linear regressor
```

```
In [12]: plt.plot(range(epochs), (loss_record))
```


```
Out[12]: [<matplotlib.lines.Line2D at 0x7fcd542e4610>]
```



```
In [13]: xs = np.linspace(0,10,200)
xs = xs[:,np.newaxis]
xsP, _ = ml.transforms.rescale(ml.transforms.fpoly(xs,degree=degree,bias=False), params)
xsP_tensor = torch.from_numpy(xsP).float()
ys = linear_regressor(xsP_tensor)

plt.scatter(Xtr,Ytr,label="training data")
plt.plot(xs,ys.detach().numpy(),label="prediction function",color = 'red')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
```

```
Out[13]: <matplotlib.legend.Legend at 0x7fcd543d5f10>
```



Statement of Collaboration

I, Andy Quoc Anh Dinh Tran, did this assignment by myself.