

Problem 1

```
In [1]: import numpy as np
        from datetime import datetime

        import torch
        import torch.nn as nn
        import torch.nn.functional as F
        from torch.utils.data import DataLoader

        from torchvision import datasets, transforms

        %matplotlib inline
        import matplotlib.pyplot as plt
```

```
In [2]: # define transforms
        transforms = transforms.Compose([transforms.Resize((32, 32)),
                                         transforms.ToTensor()])

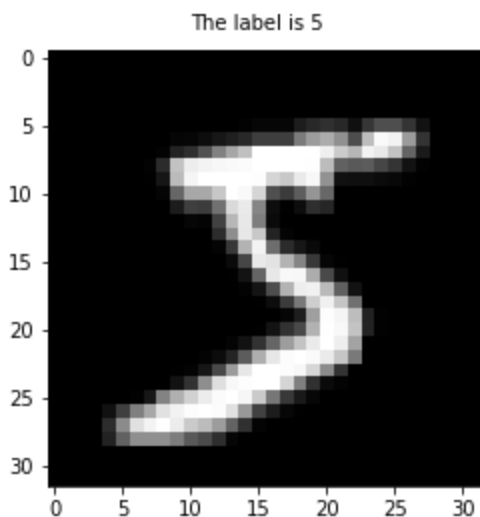
        # download and create datasets
        train_dataset = datasets.MNIST(root='mnist_data',
                                       train=True,
                                       transform=transforms,
                                       download=True)

        valid_dataset = datasets.MNIST(root='mnist_data',
                                       train=False,
                                       transform=transforms)
```

1.1.1

```
In [3]: plt.imshow(train_dataset[0][0][0], cmap='gray')
        plt.text(10, -2, 'The label is ' + str("5"))
```

Out[3]: Text(10, -2, 'The label is 5')



```
In [4]: # hyper parameters
        RANDOM_SEED = 42
        LEARNING_RATE = 0.001
        BATCH_SIZE = 32
        N_EPOCHS = 15
```

```
IMG_SIZE = 32
N_CLASSES = 10
```

1.1.2

```
In [5]: # define the data loaders
train_loader = DataLoader(dataset=train_dataset,
                           batch_size=BATCH_SIZE,
                           shuffle=True)

valid_loader = DataLoader(dataset=valid_dataset,
                           batch_size=BATCH_SIZE,
                           shuffle=True)
```

1.1.3

```
In [6]: def train(train_loader, model, criterion, optimizer):
        '''
        Train one epoch.
        '''

        model.train()
        running_loss = 0

        for X, y_true in train_loader:

            optimizer.zero_grad()

            # Forward pass
            y_hat, _ = model(X)
            loss = criterion(y_hat, y_true)
            running_loss += loss.item() * X.size(0)

            # Backward pass
            loss.backward()
            optimizer.step()

        epoch_loss = running_loss / len(train_loader.dataset)
        return model, optimizer, epoch_loss
```

1.1.4

```
In [7]: def validate(valid_loader, model, criterion):
        '''
        Function for the validation step of the training loop.
        Returns the model and the loss on the test set.
        '''

        model.eval()
        running_loss = 0

        for X, y_true in valid_loader:

            # Forward pass and record loss
            y_hat, _ = model(X)
            loss = criterion(y_hat, y_true)

            running_loss += loss.item() * X.size(0)
```

```
epoch_loss = running_loss / len(valid_loader.dataset)

return model, epoch_loss
```

```
In [8]: def training_loop(model, criterion, optimizer, train_loader, valid_loader, epochs, print_every):
'''
Function defining the entire training loop
'''

# set objects for storing metrics
best_loss = 1e10
train_losses = []
valid_losses = []
train_accs = []
valid_accs = []

# Train model
for epoch in range(0, epochs):

    # training
    model, optimizer, train_loss = train(train_loader, model, criterion, optimizer)
    train_losses.append(train_loss)

    # validation
    with torch.no_grad():
        model, valid_loss = validate(valid_loader, model, criterion)
        valid_losses.append(valid_loss)

    if epoch % print_every == (print_every - 1):

        train_acc = get_accuracy(model, train_loader,)
        train_accs.append(train_acc)
        valid_acc = get_accuracy(model, valid_loader)
        valid_accs.append(valid_acc)

        print(f'{datetime.now().time().replace(microsecond=0)} '
              f'Epoch: {epoch}\t'
              f'Train loss: {train_loss:.4f}\t'
              f'Valid loss: {valid_loss:.4f}\t'
              f'Train accuracy: {100 * train_acc:.2f}\t'
              f'Valid accuracy: {100 * valid_acc:.2f}')

    performance = {
        'train_losses': train_losses,
        'valid_losses': valid_losses,
        'train_acc': train_accs,
        'valid_acc': valid_accs
    }

    return model, optimizer, performance
```

1.1.5

```
In [9]: def get_accuracy(model, data_loader):
'''
Function for computing the accuracy of the predictions over the entire data_loader
'''

correct_pred = 0
n = 0

with torch.no_grad():
```

```

model.eval()
for X, y_true in data_loader:

    y_hat, _ = model(X)
    _, predicted_labels = torch.max(y_hat.data, 1)

    n += y_true.size(0)
    correct_pred += (predicted_labels == y_true).sum()

return correct_pred.float() / n

def plot_performance(performance):
    """
    Function for plotting training and validation losses
    """

    # temporarily change the style of the plots to seaborn
    plt.style.use('seaborn')

    fig, ax = plt.subplots(1, 2, figsize = (16, 4.5))
    for key, value in performance.items():
        if 'loss' in key:
            ax[0].plot(value, label=key)
        else:
            ax[1].plot(value, label=key)
    ax[0].set(title="Loss Over Epochs",
              xlabel='Epoch',
              ylabel='Loss')
    ax[1].set(title="Accuracy Over Epochs",
              xlabel='Epoch',
              ylabel='Loss')
    ax[0].legend()
    ax[1].legend()
    plt.show()

    # change the plot style to default
    plt.style.use('default')

```

1.2.1

In [10]: `class LeNet5(nn.Module):`

```

def __init__(self, n_classes):
    super(LeNet5, self).__init__()

    self.convolution_layer = nn.Sequential(
        nn.Conv2d(in_channels=1, out_channels=6, kernel_size=5, stride=1),
        nn.Tanh(),
        nn.AvgPool2d(kernel_size=2, stride=2, padding=0),
        nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5, stride=1),
        nn.Tanh(),
        nn.AvgPool2d(kernel_size=2, stride=2, padding=0),
        nn.Conv2d(in_channels=16, out_channels=120, kernel_size=5, stride=1),
        nn.Tanh()
    )

    self.linear_layer = nn.Sequential(
        nn.Linear(in_features=120, out_features=84),
        nn.Tanh(),
        nn.Linear(in_features=84, out_features=n_classes)
    )

```

```

def forward(self, x):
    x = self.convolution_layer(x)
    x = torch.flatten(x, 1)
    x = self.linear_layer(x)
    #logits = self.classifier(x)
    logits = x
    probs = F.softmax(logits, dim=1)
    return logits, probs

```

1.2.2

In [11]: `class MLP(nn.Module):`

```

def __init__(self, layers):
    super(MLP, self).__init__()

    self.all_layers = []
    self.all_layers.append(nn.Flatten())
    self.all_layers.append(nn.Linear(layers[0], layers[1]))
    for i in range(len(layers) - 2):
        self.all_layers.append(nn.Tanh())
        self.all_layers.append(nn.Linear(layers[i+1], layers[i+2]))

    self.layers = nn.Sequential(*self.all_layers)

def forward(self, x):
    x = self.layers(x)
    logits = x
    probs = F.softmax(logits, dim=1)
    return logits, probs

```

1.3.1

In [12]: `torch.manual_seed(RANDOM_SEED)`

```

model = LeNet5(N_CLASSES)
print(model)
optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)
criterion = nn.CrossEntropyLoss()

```

```


LeNet5(
  (convolution_layer): Sequential(
    (0): Conv2d(1, 6, kernel_size=(5, 5), stride=(1, 1))
    (1): Tanh()
    (2): AvgPool2d(kernel_size=2, stride=2, padding=0)
    (3): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
    (4): Tanh()
    (5): AvgPool2d(kernel_size=2, stride=2, padding=0)
    (6): Conv2d(16, 120, kernel_size=(5, 5), stride=(1, 1))
    (7): Tanh()
  )
  (linear_layer): Sequential(
    (0): Linear(in_features=120, out_features=84, bias=True)
    (1): Tanh()
    (2): Linear(in_features=84, out_features=10, bias=True)
  )
)

```

```
In [13]: model, optimizer, performance_1 = training_loop(model, criterion, optimizer, train_loader,
#plot_performance(performance_1)
```

Time	Epoch	Train loss	Valid loss	Train accuracy
18:12:51	0	0.2290	0.1020	96.84
				Valid accuracy: 96.81
18:13:22	1	0.0766	0.0681	98.20
				Valid accuracy: 97.95
18:13:53	2	0.0538	0.0573	98.60
				Valid accuracy: 98.25
18:14:25	3	0.0432	0.0473	99.03
				Valid accuracy: 98.53
18:14:56	4	0.0346	0.0529	99.08
				Valid accuracy: 98.30
18:15:27	5	0.0306	0.0493	99.20
				Valid accuracy: 98.44
18:15:58	6	0.0253	0.0409	99.40
				Valid accuracy: 98.76
18:16:29	7	0.0216	0.0446	99.58
				Valid accuracy: 98.64
18:17:00	8	0.0194	0.0444	99.49
				Valid accuracy: 98.72
18:17:31	9	0.0187	0.0450	99.52
				Valid accuracy: 98.68
18:18:02	10	0.0150	0.0595	99.35
				Valid accuracy: 98.36
18:18:33	11	0.0145	0.0435	99.69
				Valid accuracy: 98.85
18:19:04	12	0.0129	0.0613	99.28
				Valid accuracy: 98.24
18:19:35	13	0.0142	0.0363	99.75
				Valid accuracy: 98.95
18:20:06	14	0.0101	0.0458	99.63
				Valid accuracy: 98.82

```
In [14]: plot_performance(performance_1)
```



1.3.2

```
In [15]: torch.manual_seed(RANDOM_SEED)
```

```
layers = [1024, 256, 64, 16, N_CLASSES]
model = MLP(layers)
print(model)
optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)
criterion = nn.CrossEntropyLoss()
```

```
MLP(
  (layers): Sequential(
    (0): Flatten(start_dim=1, end_dim=-1)
```

```

(1): Linear(in_features=1024, out_features=256, bias=True)
(2): Tanh()
(3): Linear(in_features=256, out_features=64, bias=True)
(4): Tanh()
(5): Linear(in_features=64, out_features=16, bias=True)
(6): Tanh()
(7): Linear(in_features=16, out_features=10, bias=True)
)
)

```

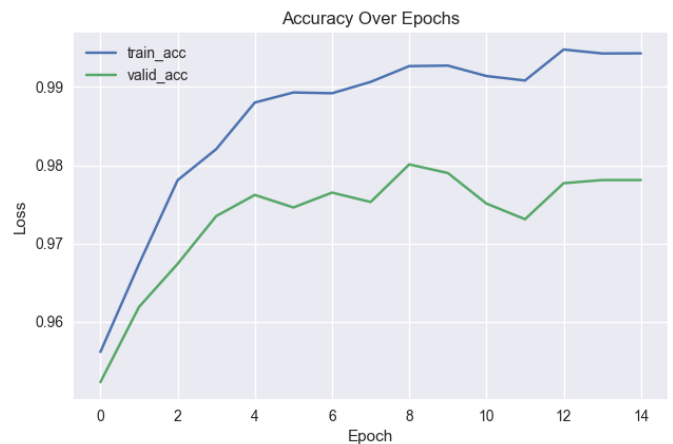
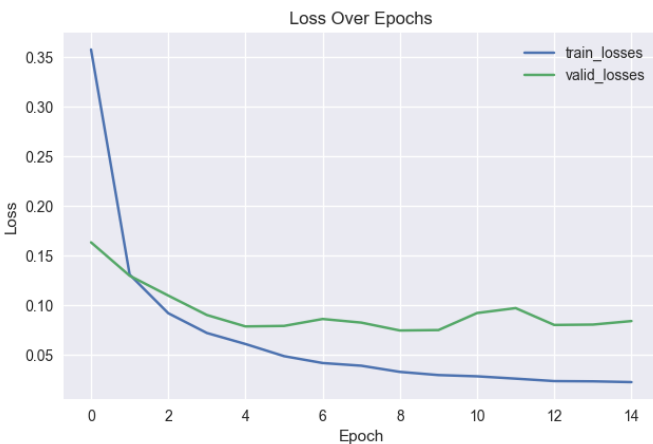
```
In [16]: model, optimizer, performance_2 = training_loop(model, criterion, optimizer, train_loader,
```

```

18:20:27 Epoch: 0      Train loss: 0.3575      Valid loss: 0.1636      Train accuracy:
95.61   Valid accuracy: 95.23
18:20:48 Epoch: 1      Train loss: 0.1311      Valid loss: 0.1300      Train accuracy:
96.74   Valid accuracy: 96.19
18:21:09 Epoch: 2      Train loss: 0.0923      Valid loss: 0.1100      Train accuracy:
97.81   Valid accuracy: 96.74
18:21:30 Epoch: 3      Train loss: 0.0724      Valid loss: 0.0905      Train accuracy:
98.21   Valid accuracy: 97.35
18:21:51 Epoch: 4      Train loss: 0.0613      Valid loss: 0.0790      Train accuracy:
98.80   Valid accuracy: 97.62
18:22:12 Epoch: 5      Train loss: 0.0491      Valid loss: 0.0795      Train accuracy:
98.93   Valid accuracy: 97.46
18:22:33 Epoch: 6      Train loss: 0.0421      Valid loss: 0.0864      Train accuracy:
98.92   Valid accuracy: 97.65
18:22:54 Epoch: 7      Train loss: 0.0394      Valid loss: 0.0828      Train accuracy:
99.06   Valid accuracy: 97.53
18:23:15 Epoch: 8      Train loss: 0.0332      Valid loss: 0.0749      Train accuracy:
99.26   Valid accuracy: 98.01
18:23:36 Epoch: 9      Train loss: 0.0300      Valid loss: 0.0753      Train accuracy:
99.27   Valid accuracy: 97.90
18:23:56 Epoch: 10     Train loss: 0.0288      Valid loss: 0.0924      Train accuracy:
99.14   Valid accuracy: 97.51
18:24:17 Epoch: 11     Train loss: 0.0264      Valid loss: 0.0974      Train accuracy:
99.08   Valid accuracy: 97.31
18:24:38 Epoch: 12     Train loss: 0.0240      Valid loss: 0.0804      Train accuracy:
99.48   Valid accuracy: 97.77
18:24:59 Epoch: 13     Train loss: 0.0237      Valid loss: 0.0808      Train accuracy:
99.43   Valid accuracy: 97.81
18:25:20 Epoch: 14     Train loss: 0.0229      Valid loss: 0.0844      Train accuracy:
99.43   Valid accuracy: 97.81

```

```
In [17]: plot_performance(performance_2)
```



```
In [18]: #performance
```

1.4.1

The total number of trainable parameters of LeNet is 61750.

Layer	# Filters / Neurons	Filter Size	Stride	Size of Feature Map	Acitvation Function	Bias Terms	Trainable Parameter Formula	Trainable Parameters
Input	-	-	-	32 x 32 x 1	-	-	-	-
Conv 1	6	5 * 5	1	28 x 28 x 6	tanh	1	$((5 \cdot 5 \cdot 1) + 1) \cdot 6$	156
Avg. Pooling 1		2 * 2	2	14 x 14 x 6	-	1	$6 \cdot 2$	12
Conv 2	16	5 * 5	1	10 x 10 x 16	tanh	1	$((5 \cdot 5 \cdot 6) + 1) \cdot 16$	2416
Avg. Pooling 2		2 * 2	2	5 x 5 x 16	-	1	$16 \cdot 2$	32
Conv 3	120	5 * 5	1	120	tanh	1	$((5 \cdot 5 \cdot 16) + 1) \cdot 120$	48120
Fully Connected 1	-	-	-	84	tanh	1	$(120 + 1) \cdot 84$	10164
Fully Connected 2	-	-	-	10	Softmax	1	$(84 + 1) \cdot 10$	850

1.4.2

The total number of trainable parameters of MLP is 280058.

Layer	Input Size	Output Size	Acitvation Function	Bias Terms	Trainable Parameter Formula	Trainable Parameters
Input Layer	-	1024	-	-	-	-
Hidden Layer 1	1024	256	tanh	1	$(1024 + 1) \cdot 256$	262400
Hidden Layer 2	256	64	tanh	1	$(256 + 1) \cdot 64$	16448
Hidden Layer 3	64	16	tanh	1	$(64 + 1) \cdot 16$	1040
Output Layer	16	10	tanh	1	$(16 + 1) \cdot 10$	170

1.4.3

The model that has better performance in terms of prediction accuracy on the test data is LeNet and this could be because LeNet is a Convolutional Neural Network and has layers that are not fully connected comaped to Multi Layer Perceptron.

Statement of Collaboration

I, Andy Quoc Anh Dinh Tran, did this assignment by myself.